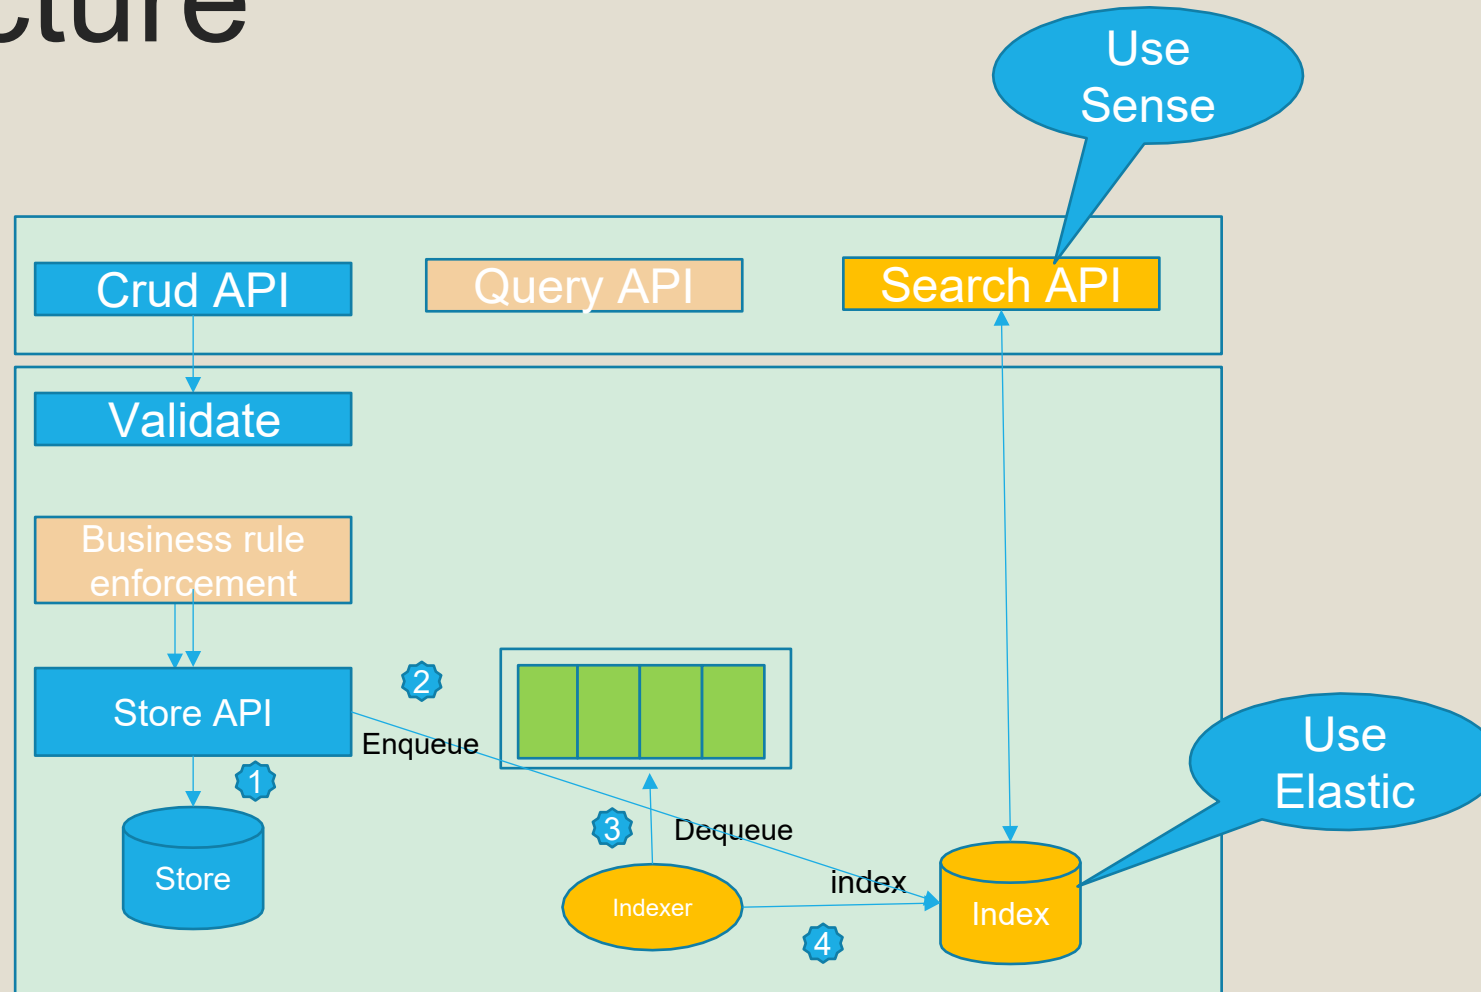




# ADVANCED TOPIC IN BIG DATA

10/22/16

# Architecture



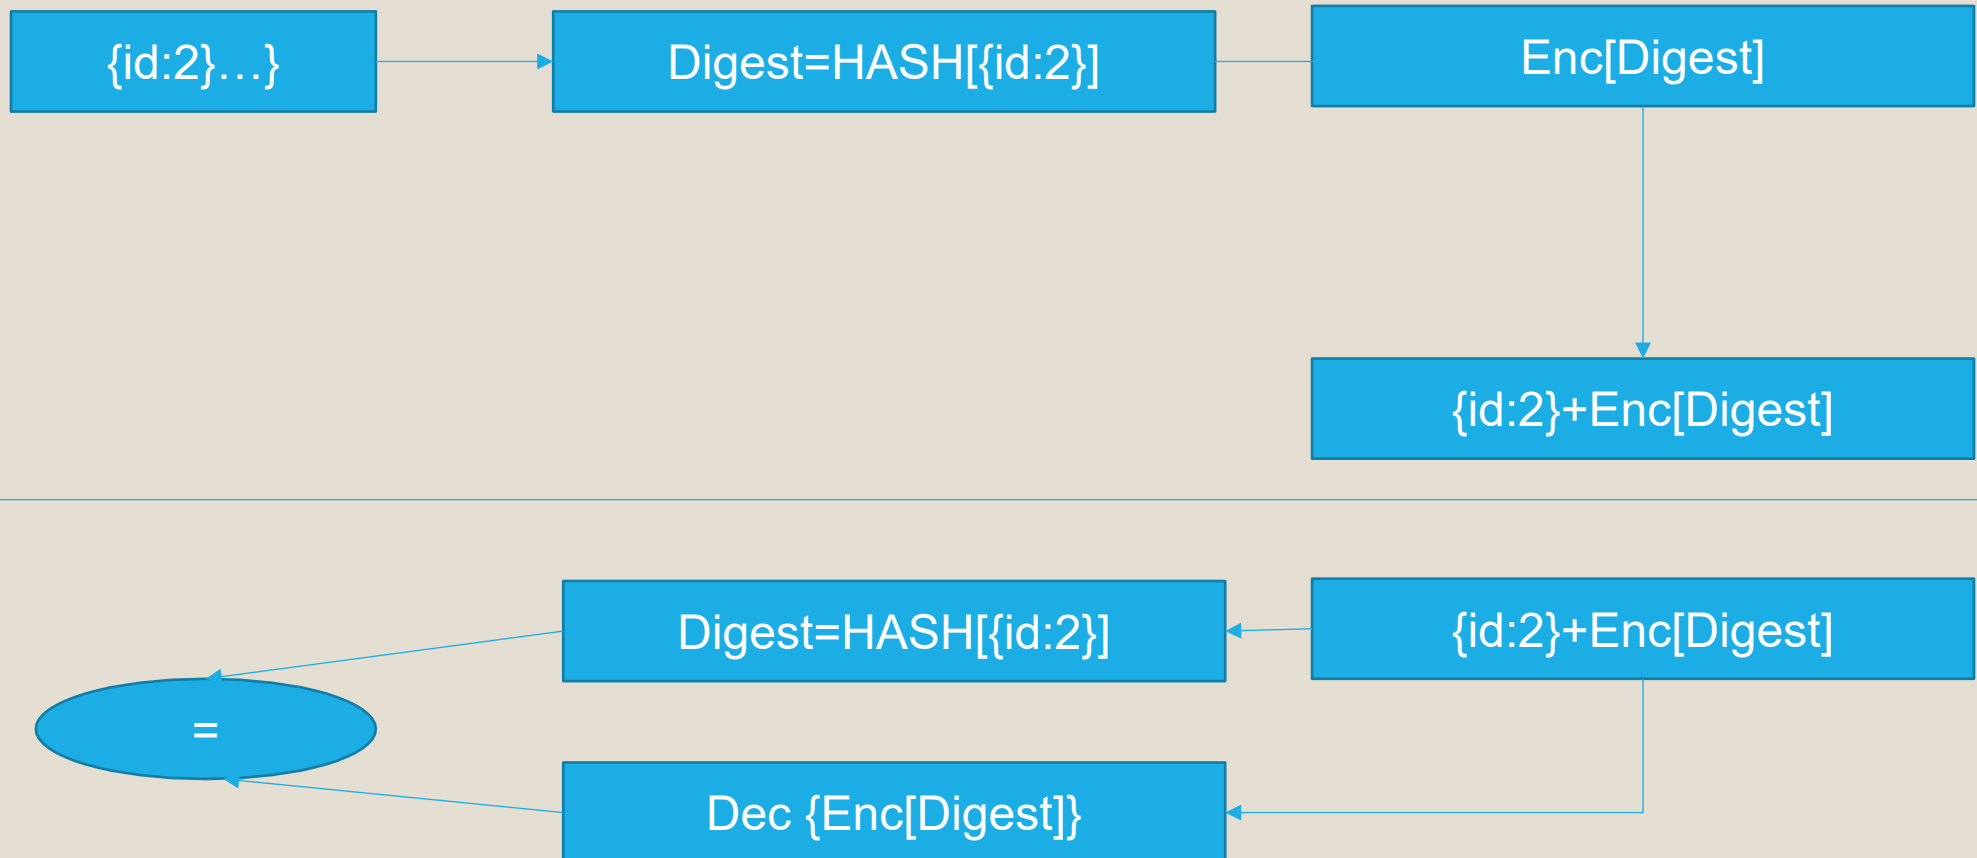
# Prototype outline:

- Rest API that can handle any structured data in Json
- Rest API with support for crud operations, including merge support, cascaded delete
- Rest API with support for validation
- Json Schema describing the data model for the use case
- Advanced semantics with rest API operations such as update if not changed
- Storage of data in key/value store
- Search with join using Elastic
  - Parent-Child indexing
- Queueing
- Security

# Catching up on old topics

- Implementing merge semantics
- Etag implementations
- Security implementation progress
  - Revisit security implementation using tokens

# Asymmetric Crypto



# Join and elastic search

- How to implement to join with elastic search?
  - Pros
  - Cons
- Demonstrate this approach using sense

<http://fideloper.com/api-etag-conditional-get>

They use MD5 having to calculate the Etag.

<http://stackoverflow.com/questions/415953/how-can-i-generate-an-md5-hash>

# Fulltext search

- Basic concepts:

- Indexing:

- Is the process of creating an index.
    - An index is defined as a collection of fields. Each field can be either single value/multivalued, have a type, stored, indexed, required, can be associated with different tokenizer's/analyzers
    - Dynamic fields is a very useful feature
    - An index contains a collection of documents.
    - A document is a collection of property (field) / value pairs

- Searching

- Is the process of discovering a document in an index that meets certain criteria's
    - the criterias are specified using fields that are found in a document



# Query samples

- find all documents containing name:jeff
- find all documents containing name:jeff and age:30 (Or any other logical relation, e.g. or, not, and)
- find all documents created after 9-16-2016
- find all documents of type plan
- find all documents of type pla\*; E.g.; plans, planning, planner etc.
- Find all the Unique terms of the field "type" in the system
- Counts:
  - how many times a certain value occurs in the index
- Aggregates:
  - Max, Min, Average, Sum, percentiles, etc.
- How many cameras are on sale between 50 and \$100?

# Faceted queries

- Is the bucketing of search results into buckets based on terms in the index
- Useful for determining the unique terms for a field and returns a count for each of those terms.
- Makes it easy to explore search results
- Faceting example is found here:
  - <https://lucidworks.com/post/faceted-search-with-solr/>

# Faceting..

- Field faceting – retrieve the counts for all terms, or just the top terms in any given field. The field must be indexed.
- Query faceting – return the number of documents in the current search results that also match the given query.
- Date faceting – return the number of documents that fall within certain date ranges.

# Filter queries

- Used to filter the results of the previous query
  - Often used to implement drill down into search results
- When filter query is added to the previous query, its effect is to exclude results that do not match the filter
- Example:
  - Return all cameras by manufacturer and their count
    - `/query?q=camera facet.field=manu`
  - Return all cameras in this price range by manufacturer and their count
    - `http://localhost:8983/solr/query?q=camera &facet.field=manu &fq=price:[400 to 500] (fq is filter query)`

# Elastic Search

- Getting started:
  - <https://www.elastic.co/guide/en/elasticsearch/guide/current/getting-started.html>

# Homework

- Demonstrate an example of join queries using elastic search. This is due 10/22

# References

- <https://cwiki.apache.org/confluence/display/solr/About+This+Guide>
- <https://lucidworks.com/post/faceted-search-with-solr/https://www.elastic.co/guide/en/elasticsearch/guide/current/denormalization.html>
- Getting started:
  - <https://www.elastic.co/guide/en/elasticsearch/guide/current/getting-started.html>

# Oauth

INFO 7255



# Use cases for security

- Washington Post/Boston Globe: paywall, tiered-based subscription
- Flash sales
  - Can I prevent bots from sweeping up all inventory?
  - Can my application hold up against excessive demand?
    - Digital Waiting Room
- Authenticated access
  - Quota and throttling
- Anonymous access
  - throttling
- Bots Access
  - Good bots versus bad bots

# Security requirements

- Authorized access against API
  - Only users authorized to access resources are allowed
  - Users able to see/edit their own plans
  - Users may read other plans, but no change them
    - Users may have certain access to this endpoint but not to the other one
- Anonymous browsing may be allowed
  - This is prior to user authentication
- App may not exceed certain requests per day/month: quota
- Apps that are making excessive number of requests need to be throttled
  - Digital Waiting Room

# High-level Approach

- Client includes an authorization header
  - The value of the header is a token
- API uses the authorization header value (token) for authorization and authentication
  - Client signs token
  - API verifies token

# Key design questions

- What is the overall approach for securing APIs?
  - Bearer Tokens
- What is the token structure?
  - JWT
- How are token generated?
  - How are they signed?
    - By an Idp
- How are tokens verified?
  - Authenticate the signer of the token
- Security crypto: Asymmetric? RS256
- Security guarantees: Authentication, non-tampering

# First approach: API keys for securing access by apps

- High level flow:
  - Each app is granted a key at build time by the server
  - app includes key in every request that goes to server
  - Implications on quota and throttling?

# OAUTH 1.0

- Username & Password

# Industry accepted approach: OAUTH 2.0

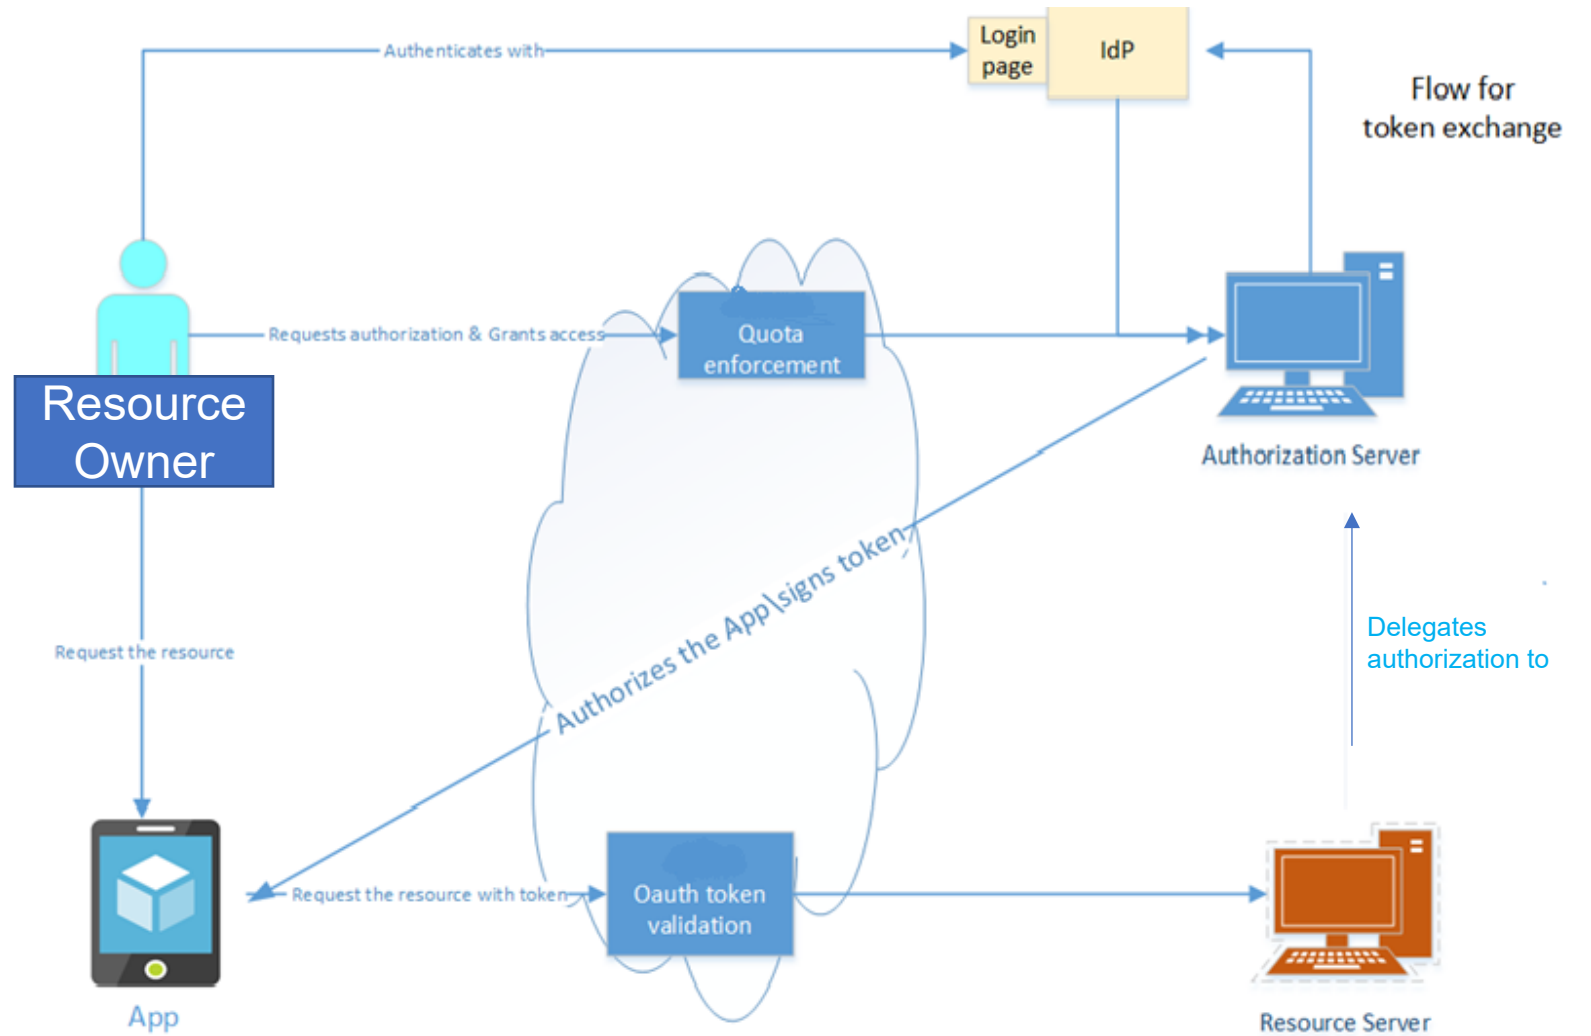
- User downloads an app
- User authenticates with an IDP/Auth server
- User consents to give app access to user's data
- IDP generates token
- App includes the token in the API calls

# Public versus private app

- Public apps are those that cannot secure their credentials: single page application, mobile apps
- Private apps are those that can secure their credentials: any app running behind a firewall



# oAUTH 2.0 Overview and Actors



# Token Validation by Resource Server

- 1. Validate the structure of a JWT
- 2. Create an “allow list” that contains valid values for iss claim
- 3. Base64decode JWT header, payload
- 4. Retrieve alg and kid from Header
- 5. Retrieve iss from payload
- 6. Compare the value of iss to that stored in the “allow list”
  - 5. If iss value in allow list, use JWKS\_URI to retrieve public key.  
Otherwise, signature invalid
  - 6. Verify signature
  - 7. Validate any other claims such as scope, aud, exp, etc.

# Overview

- RFC OAUTH 2.0: <https://tools.ietf.org/html/rfc6749>
- JWT <https://tools.ietf.org/html/rfc7519>
- Example: <https://dev.fitbit.com/docs/oauth2/>

# Oauth provider (Authorization Server)

- /register
- /Authorize
  - unsecure
  - Authorization code grant flow
    - Returns both access token and refresh token
    - Use for secure clients
  - Authorization code grant flow with PKCE
    - Use for unsecured client
  - implicit grant flow
    - Returns only access token
    - Use for unsecured client
- /Token
  - Secure
  - Exchange authorization code for a token
  - generate a new token from a refresh token

# /register

- Input:
  - Client\_type = confidential (private) or public
  - redirect\_URI: https://
- Output:
  - client\_ID, client\_secret if client is confidential
  - Client\_id for public

# /Authorize

- The authorization endpoint must support "get"
- The supported query parameters are:
  - response\_type  
REQUIRED. Value MUST be either "code" or "token"
  - client\_id  
REQUIRED. The client identifier obtained from the registration
  - redirect\_uri  
Required. As described in Section 3.1.2.  
https
  - scope  
Required.
  - state  
Required

# Authorization grant code flow example:

GET /authorize?response\_type=code&client\_id=s6BhdRkqt3&state=xyz  
&redirect\_uri=https%3A%2F%2Fclient%2Eexample%2Ecom%2Fcb &scope=read  
HTTP/1.1

Host: [server.example.com](https://server.example.com)

- The response should have:
- HTTP status code should be set to 302:
- the redirect URI as the value of the location header.
- the query parameter code and its value, state and its value appended to the redirect URI.
- The code is required at all times. The state is required only if it has been present in the request.
- example:
- HTTP/1.1 302 Found
- Location: and that we should specify the location and the location it should contain it [it&state=xyz](https://server.example.com/authorize?code=123&state=xyz)

# error

- HTTP/1.1 302 Found
  - Location: [https://client.example.com/cb?error=access\\_denied&state=xyz](https://client.example.com/cb?error=access_denied&state=xyz)
  - REQUIRED. A single ASCII [USASCII] error code from the following:
    - invalid\_request
      - The request is missing a required parameter, includes an
      - invalid parameter value, includes a parameter more than
      - once, or is otherwise malformed.
    - unauthorized\_
      - The client is not authorized to request an authorization
      - code using this method.
    - access\_denied
      - The resource owner or authorization server denied the request.
    - unsupported\_response\_type
      - The authorization server does not support obtaining an
      - authorization code using this method.
    - invalid\_scope
      - The requested scope is invalid, unknown, or malformed.
    - server\_error
      - The authorization server encountered an unexpected condition that prevented it from fulfilling the request.
  - temporarily\_unavailable
    - The authorization server is currently unable to handle the request due to a temporary overloading or maintenance of the server. (This error code is needed because a 503 Service Unavailable HTTP status code cannot be returned to the client via an HTTP redirect.)
- error\_description
- OPTIONAL. Human-readable ASCII [USASCII] text providing additional information, used to assist the client developer in understanding the error that occurred. Values for the "error\_description" parameter MUST NOT include characters outside the set %x20-21 / %x23-5B / %x5D-7E.
- error\_uri
- OPTIONAL. A URI identifying a human-readable web page with information about the error, used to provide the client developer with additional information about the error. Values for the "error\_uri" parameter MUST conform to the URI-reference syntax and thus MUST NOT include characters outside the set %x21 / %x23-5B / %x5D-7E.
- state
- REQUIRED if a "state" parameter was present in the client authorization request. The exact value received from the client.



# Implicit grant flow

- GET  
/authorize?response\_type=token&client\_id=s6BhdRkqt3&state=xyz&redirect\_uri=https%3A%2F%2Fclient%2Eexample%2Ecom%2Fcb&state=xyz&scope=read
- Host: [server.example.com](https://server.example.com)
- the authorization server issues an access token and delivers it to the client by adding
- the following parameters to the fragment component of the redirectionURI:
- HTTP/1.1 302 Found
- Location: [https://example.com/cb#access\\_token=2YotnFZFEjr1zCsicMWpAA&state=xyz&token\\_type=Bearer&expires\\_in=3600](https://example.com/cb#access_token=2YotnFZFEjr1zCsicMWpAA&state=xyz&token_type=Bearer&expires_in=3600)
- access\_token
- REQUIRED. The access token issued by the authorization server.
- token\_type
- REQUIRED. The value should be set to bearer
- expires\_in
- RECOMMENDED. The lifetime in seconds of the access token
- scope
- REQUIRED, if identical to the scope requested by the client;
- otherwise, REQUIRED. The scope of the access token as described by Section 3.3.
- state
- REQUIRED if the "state" parameter was present in the client authorization request. The exact value received from the client.

# HTTP/1.1 302 Found

Location: https://client.example.com/cb#error=access\_denied&state=xy

- error
- REQUIRED. A single ASCII [USASCII] error code from the following:
- invalid\_request
- The request is missing a required parameter, includes an invalid parameter value, includes a parameter more than

once, or is otherwise malformed.

- unauthorized\_client
- The client is not authorized to request an access token using this method.
- access\_denied
- The resource owner or authorization server denied the request.
- unsupported\_response\_type
- The authorization server does not support obtaining an access token using this method.
- invalid\_scope
- The requested scope is invalid, unknown, or malformed.
- server\_error
- The authorization server encountered an unexpected condition that prevented it from fulfilling the request.
- (This error code is needed because a 500 Internal Server Error HTTP status code cannot be returned to the client
- via an HTTP redirect.)

- temporarily\_unavailable
- The authorization server is currently unable to handle the request due to a temporary overloading or maintenance
- of the server. (This error code is needed because a 503 Service Unavailable HTTP status code cannot be returned
- to the client via an HTTP redirect.)
- Values for the "error" parameter MUST NOT include characters outside the set %x20-21 / %x23-5B / %x5D-7E.
- error\_description
- OPTIONAL. Human-readable ASCII [USASCII] text providing additional information, used to assist the client developer in
- understanding the error that occurred. Values for the "error\_description" parameter MUST NOT include
- characters outside the set %x20-21 / %x23-5B / %x5D-7E.
- error\_uri
- OPTIONAL. A URI identifying a human-readable web page with
- information about the error, used to provide the client developer with additional information about the error.
- Values for the "error\_uri" parameter MUST conform to the URI-reference syntax and thus MUST NOT include characters
- outside the set %x21 / %x23-5B / %x5D-7E.
- state
- REQUIRED if a "state" parameter was present in the client authorization request. The exact value received from the
- client

# /token

- The token request endpoint must support post with Content-Type: application/x-www-form-urlencoded
- the token request endpoint must authenticate the client making the request
- the token request end point must support basic authentication
- the token endpoint must ensure that the authorization code was issued to this client\_ID
- the token endpoint must ensure that the authorization code is valid.
- the token endpoint must ensure that the authorization code is used ONLY once.
- authorization code must expire in 10s of seconds
- The token endpoint must set Cache-Control: no-store, Pragma: no-cache headers
- The token request endpoint supports the following parameters:
  - grant\_type with value set to authorization\_code, client\_credentials, password, or refresh\_token
  - code with its value set to the authorization code
  - redirect\_URI with its value set to the redirect URI that was provided in the request for the authorization code
  - client\_id; this value is required if the client is not authenticating with the authorization server
- The return payload must include the following:  
access\_token, token\_type, expires\_in refresh\_token, and any other key value pairs.

# Exchange an authorization code for a token

- POST /token HTTP/1.1  
Host: [server.example.com](https://server.example.com)  
Authorization: Basic czZCaGRSa3F0MzpnWDFmQmF0M2JW  
Content-Type: application/x-www-form-urlencoded
- grant\_type=authorization\_code&code=SpIxlOBzQqYbYS6WxSbIA  
&redirect\_uri=https%3A%2F%2Fclient%2Eexample%2Ecom%2Fcb
- HTTP/1.1 200 OK  
Content-Type: application/json; charset=UTF-8  
Cache-Control: no-store  
Pragma: no-cache

```
{  
  "access_token": "2YotnFZFEjr1zCsicMWpAA", "token_type": "Bearer", "expires_in": 3600,  
  "refresh_token": "tGzv3JOkF0XG5Qx2TlKWIA",  
  "example_parameter": "example_value"  
}
```

# Refreshing the access token

- POST /token HTTP/1.1  
Host: [server.example.com](http://server.example.com)  
Authorization: Basic czZCaGRSa3F0MzpnWDFmQmF0M2JW  
Content-Type: application/x-www-form-urlencoded
- grant\_type=refresh\_token&refresh\_token=tGzv3JOkF0XG5Qx2  
TIKWIA

# Methodology for securing rest API

- Client app registers with OAuth/Authorization Server
- Client app request a token
- OAuth provider generates an access token to client APP
- Client app includes access token in every HTTP request using Authorization header
- Client app sets the Authorization header to Bearer {access token}
- The rest API validates the access token
  - What does it need to validate the token?

# Token Validation by Resource Server

- 1. Validate the structure of a JWT
- 2. Create an “allow list” that contains valid values for iss claims
- 3. Base64decode JWT header, payload
- 4. Retrieve alg and kid from Header
- 5. Retrieve iss from payload
- 6. Compare the value of iss to that stored in the “allowed list”
  - 5. If iss value in allow list, use JWKS to retrieve public key. Otherwise, signature invalid
  - 6. Verify signature
  - 7. Validate any other claims such as scope, aud, exp, etc.

# JWT example

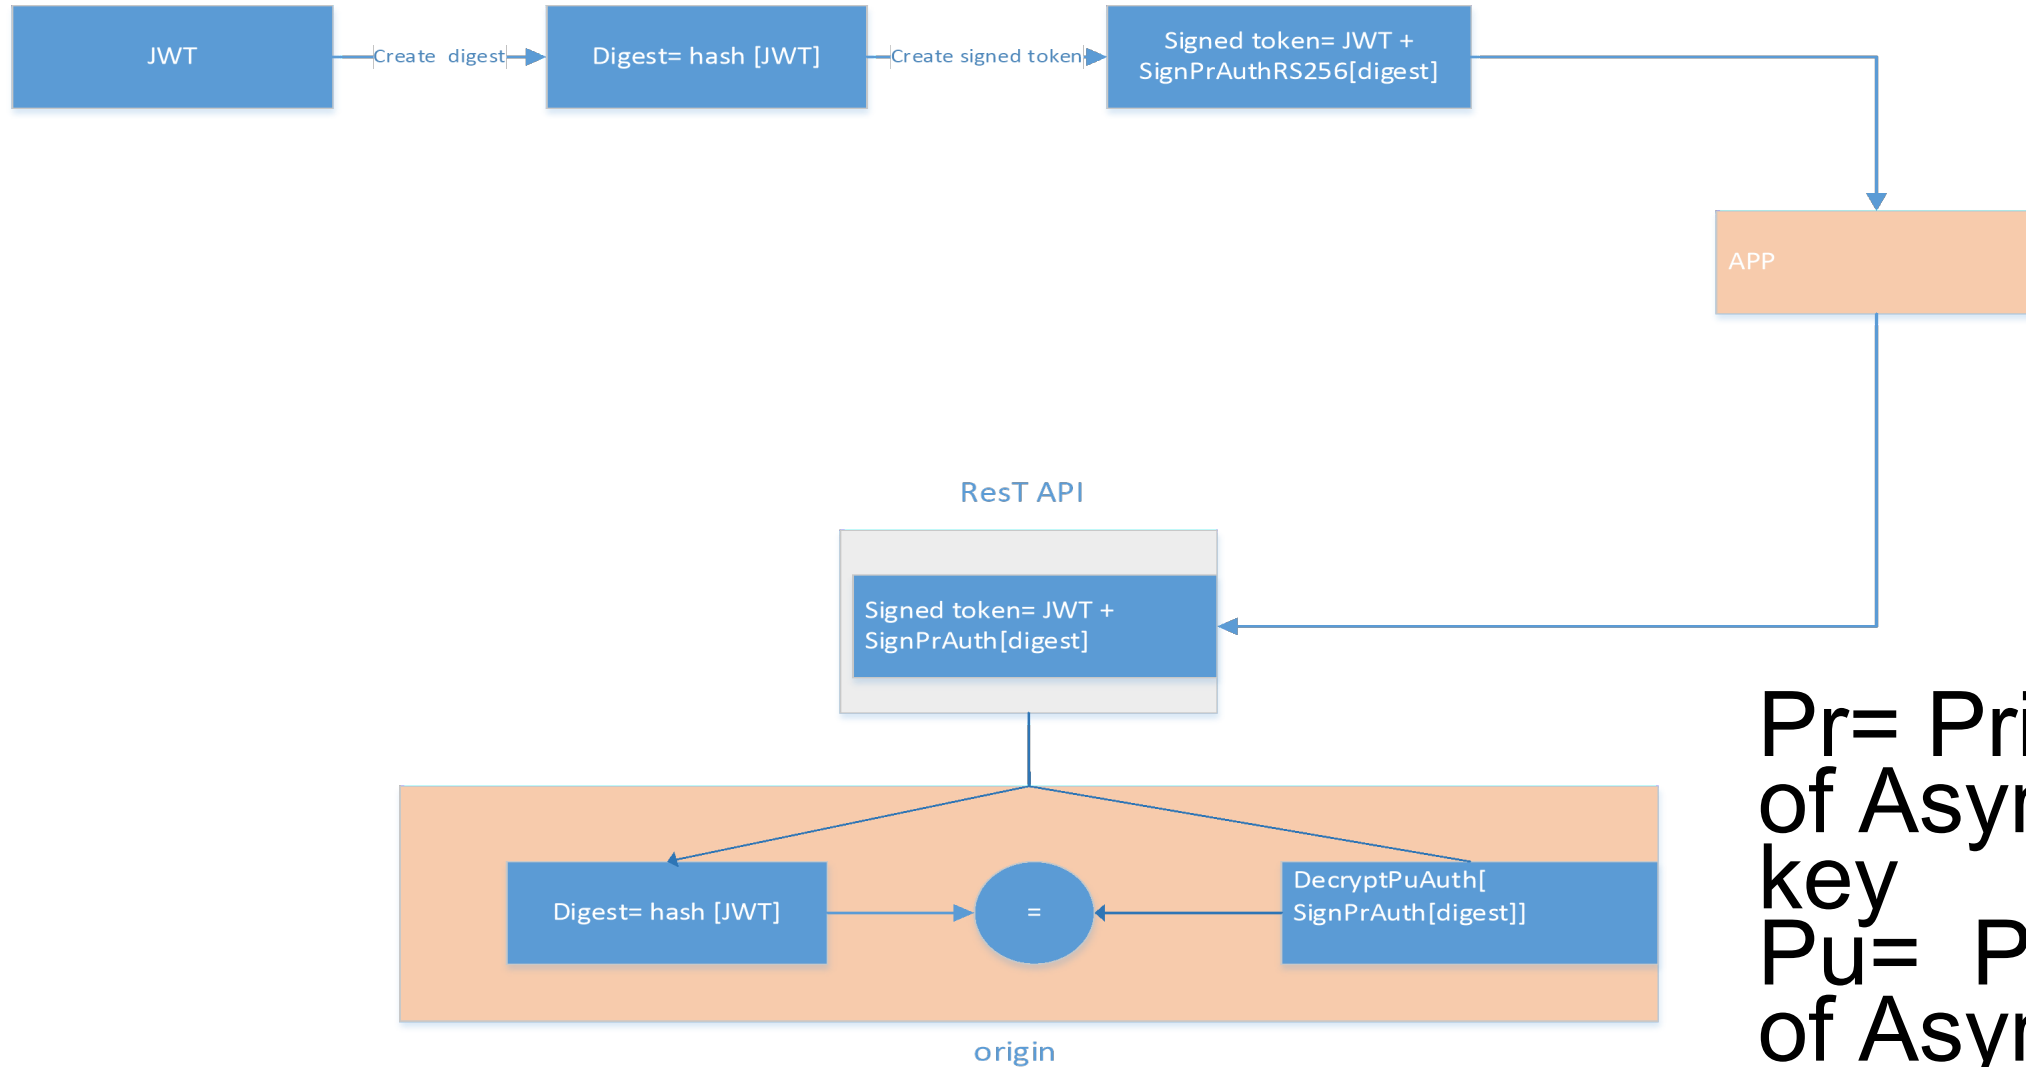
```
{  
  "typ": "JWT"  
}
```

```
{  
  "app": "TEST",  
  "acc": "7888-a9a0-4de2-be72-57775575",  
  "iss": "yyy",  
  "scope": ["read", "write"],  
  "exp": 1561939073,  
  "jti": "jhhhjhg-6cab-lkjll-8512-kjkkjk",  
  "aud": "/plan/{id}"  
}
```

```
RSASHA256( base64UrlEncode(header) + "." + base64UrlEncode(payload)
```



# Signature Verification using RS 256



Pr= Private key  
of Asymmetric  
key  
Pu= Public key  
of Asymmetric  
key

# Key Distribution

- When using RS 256:
  - Generate a public/private key pair
  - Signer uses the private key to sign the token
  - Rest API uses the public key to verify the signature
  - Rest API must have access to the public key
    - JWK : <https://tools.ietf.org/html/rfc7517>

# Key Rotation (Private) and Distribution (Public)

- Using Kid
- {
- "alg": "RS256",
- "typ": "JWT",
- "kid": "2",
- ~~• "jku": <https://myjwks> ;;; Not recommended to include this~~
- }

# References for token signing

- <https://connect2id.com/products/nimbus-jose-jwt/examples/jwt-with-rsa-signature>
- [https://en.wikipedia.org/wiki/JSON\\_Web\\_Token](https://en.wikipedia.org/wiki/JSON_Web_Token)
- <https://tools.ietf.org/html/rfc7519>
- <https://developers.google.com/oauthplayground/>
- <https://developers.google.com/identity/protocols/oauth2/openid-connect>
- <https://console.developers.google.com/apis/credentials?project=vital-invention-306022>
- <https://accounts.google.com/.well-known/openid-configuration>
- JWT.io
- <https://developers.google.com/identity/protocols/oauth2/openid-connect>

# Oauth

INFO 7255

# Use cases for security

- Washington Post/Boston Globe: paywall, tiered-based subscription
- Flash sales
  - Can I prevent bots from sweeping up all inventory?
  - Can my application hold up against excessive demand?
    - Digital Waiting Room
- Authenticated access
  - Quota and throttling
- Anonymous access
  - throttling
- Bots Access
  - Good bots versus bad bots

# Security requirements

- Authorized access against API
  - Only users authorized to access resources are allowed
  - Users able to see/edit their own plans
  - Users may read other plans, but no change them
    - Users may have certain access to this endpoint but not to the other one
- Anonymous browsing may be allowed
  - This is prior to user authentication
- App may not exceed certain requests per day/month: quota
- Apps that are making excessive number of requests need to be throttled
  - Digital Waiting Room

# High-level Approach

- Client includes an authorization header
  - The value of the header is a token
- API uses the authorization header value (token) for authorization and authentication
  - Client signs token
  - API verifies token



# Key design questions

- What is the overall approach for securing APIs?
  - Bearer Tokens
- What is the token structure?
  - JWT
- How are token generated?
  - How are they signed?
    - By an Idp
- How are tokens verified?
  - Authenticate the signer of the token
- Security crypto: Asymmetric? RS256
- Security guarantees: Authentication, non-tampering

# First approach: API keys for securing access by apps

- High level flow:
  - Each app is granted a key at build time by the server
  - app includes key in every request that goes to server
  - Implications on quota and throttling?

# OAUTH 1.0

- Username & Password

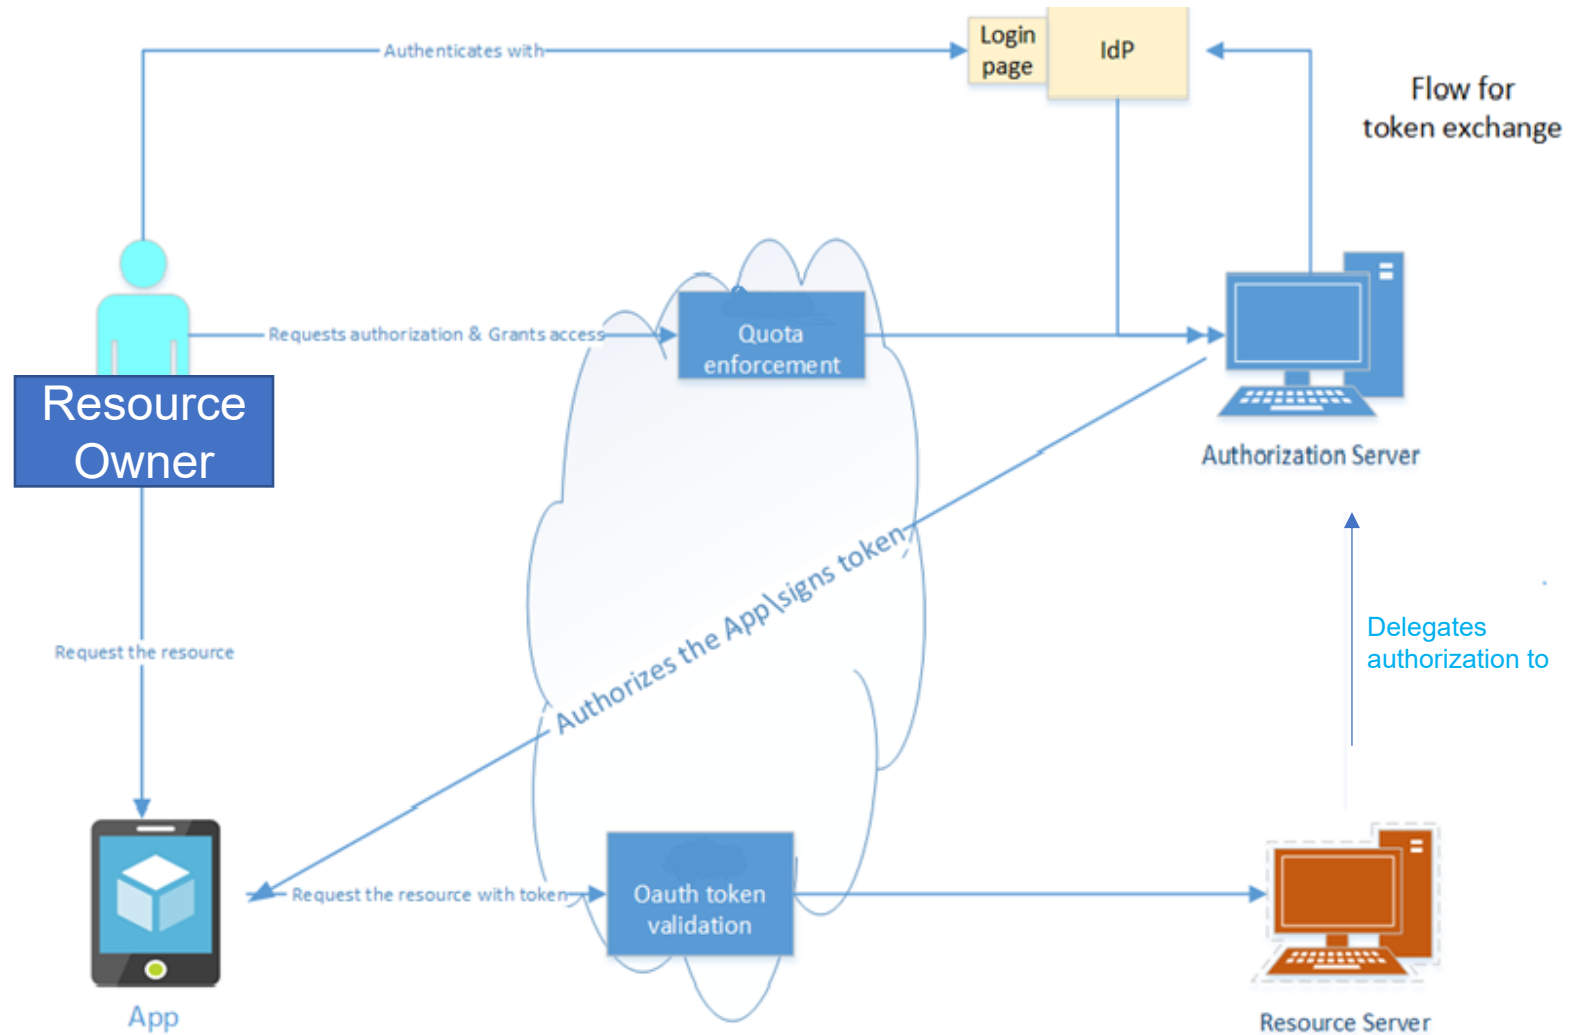
# Industry accepted approach: OAUTH 2.0

- User downloads an app
- User authenticates with an IDP/Auth server
- User consents to give app access to user's data
- IDP generates token
- App includes the token in the API calls

# Public versus private app

- Public apps are those that cannot secure their credentials: single page application, mobile apps
- Private apps are those that can secure their credentials: any app running behind a firewall

# oAUTH 2.0 Overview and Actors



# Token Validation by Resource Server

- 1. Validate the structure of a JWT
- 2. Create an “allow list” that contains valid values for iss claim
- 3. Base64decode JWT header, payload
- 4. Retrieve alg and kid from Header
- 5. Retrieve iss from payload
- 6. Compare the value of iss to that stored in the “allowed list”
  - 5. If iss value in allow list, use JWKS\_URI to retrieve public key.  
Otherwise, signature invalid
  - 6. Verify signature
  - 7. Validate any other claims such as scope, aud, exp, etc.

# Overview

- RFC OAUTH 2.0: <https://tools.ietf.org/html/rfc6749>
- JWT <https://tools.ietf.org/html/rfc7519>
- Example: <https://dev.fitbit.com/docs/oauth2/>



# Oauth provider (Authorization Server)

- /register
- /Authorize
  - unsecure
  - Authorization code grant flow
    - Returns both access token and refresh token
    - Use for secure clients
  - Authorization code grant flow with PKCE
    - Use for unsecured client
  - implicit grant flow
    - Returns only access token
    - Use for unsecured client
- /Token
  - Secure
  - Exchange authorization code for a token
  - generate a new token from a refresh token

# /register

- Input:
  - Client\_type = confidential (private) or public
  - redirect\_URI: https://
- Output:
  - client\_ID, client\_secret if client is confidential
  - Client\_id for public

# /Authorize

- The authorization endpoint must support "get"
- The supported query parameters are:
  - response\_type  
REQUIRED. Value MUST be either "code" or "token"
  - client\_id  
REQUIRED. The client identifier obtained from the registration
  - redirect\_uri  
Required. As described in Section 3.1.2.  
https
  - scope  
Required.
  - state  
Required

# Authorization grant code flow example:

GET /authorize?response\_type=code&client\_id=s6BhdRkqt3&state=xyz  
&redirect\_uri=https%3A%2F%2Fclient%2Eexample%2Ecom%2Fcb &scope=read  
HTTP/1.1

Host: [server.example.com](https://server.example.com)

- The response should have:
- HTTP status code should be set to 302:
- the redirect URI as the value of the location header.
- the query parameter code and its value, state and its value appended to the redirect URI.
- The code is required at all times. The state is required only if it has been present in the request.
- example:
- HTTP/1.1 302 Found
- Location: and that we should specify the location and the location it should contain it [it&state=xyz](https://server.example.com/authorize?code=xyz&state=xyz)

# error

- HTTP/1.1 302 Found
  - Location: [https://client.example.com/cb?error=access\\_denied&state=xyz](https://client.example.com/cb?error=access_denied&state=xyz)
  - REQUIRED. A single ASCII [USASCII] error code from the following:
    - invalid\_request
      - The request is missing a required parameter, includes an
      - invalid parameter value, includes a parameter more than
      - once, or is otherwise malformed.
    - unauthorized\_
      - The client is not authorized to request an authorization
      - code using this method.
    - access\_denied
      - The resource owner or authorization server denied the request.
    - unsupported\_response\_type
      - The authorization server does not support obtaining an
      - authorization code using this method.
    - invalid\_scope
      - The requested scope is invalid, unknown, or malformed.
    - server\_error
      - The authorization server encountered an unexpected condition that prevented it from fulfilling the request.
  - temporarily\_unavailable
    - The authorization server is currently unable to handle the request due to a temporary overloading or maintenance of the server. (This error code is needed because a 503 Service Unavailable HTTP status code cannot be returned to the client via an HTTP redirect.)
- error\_description
- OPTIONAL. Human-readable ASCII [USASCII] text providing additional information, used to assist the client developer in understanding the error that occurred. Values for the "error\_description" parameter MUST NOT include characters outside the set %x20-21 / %x23-5B / %x5D-7E.
- error\_uri
- OPTIONAL. A URI identifying a human-readable web page with information about the error, used to provide the client developer with additional information about the error. Values for the "error\_uri" parameter MUST conform to the URI-reference syntax and thus MUST NOT include characters outside the set %x21 / %x23-5B / %x5D-7E.
- state
- REQUIRED if a "state" parameter was present in the client authorization request. The exact value received from the client.

# Implicit grant flow

- GET  
/authorize?response\_type=token&client\_id=s6BhdRkqt3&state=xyz&redirect\_uri=https%3A%2F%2Fclient%2Eexample%2Ecom%2Fcb&state=xyz&scope=read
- Host: [server.example.com](https://server.example.com)
- the authorization server issues an access token and delivers it to the client by adding
- the following parameters to the fragment component of the redirectionURI:
- HTTP/1.1 302 Found
- Location: [https://example.com/cb#access\\_token=2YotnFZFEjr1zCsicMWpAA&state=xyz&token\\_type=Bearer&expires\\_in=3600](https://example.com/cb#access_token=2YotnFZFEjr1zCsicMWpAA&state=xyz&token_type=Bearer&expires_in=3600)
- access\_token
- REQUIRED. The access token issued by the authorization server.
- token\_type
- REQUIRED. The value should be set to bearer
- expires\_in
- RECOMMENDED. The lifetime in seconds of the access token
- scope
- REQUIRED, if identical to the scope requested by the client;
- otherwise, REQUIRED. The scope of the access token as described by Section 3.3.
- state
- REQUIRED if the "state" parameter was present in the client authorization request. The exact value received from the client.

# HTTP/1.1 302 Found

Location: https://client.example.com/cb#error=access\_denied&state=xy

- error
- REQUIRED. A single ASCII [USASCII] error code from the following:
- invalid\_request
- The request is missing a required parameter, includes an invalid parameter value, includes a parameter more than

once, or is otherwise malformed.

- unauthorized\_client
- The client is not authorized to request an access token using this method.
- access\_denied
- The resource owner or authorization server denied the request.
- unsupported\_response\_type
- The authorization server does not support obtaining an access token using this method.
- invalid\_scope
- The requested scope is invalid, unknown, or malformed.
- server\_error
- The authorization server encountered an unexpected condition that prevented it from fulfilling the request.
- (This error code is needed because a 500 Internal Server Error HTTP status code cannot be returned to the client
- via an HTTP redirect.)

- temporarily\_unavailable
- The authorization server is currently unable to handle the request due to a temporary overloading or maintenance
- of the server. (This error code is needed because a 503 Service Unavailable HTTP status code cannot be returned
- to the client via an HTTP redirect.)
- Values for the "error" parameter MUST NOT include characters outside the set %x20-21 / %x23-5B / %x5D-7E.
- error\_description
- OPTIONAL. Human-readable ASCII [USASCII] text providing additional information, used to assist the client developer in
- understanding the error that occurred. Values for the "error\_description" parameter MUST NOT include
- characters outside the set %x20-21 / %x23-5B / %x5D-7E.
- error\_uri
- OPTIONAL. A URI identifying a human-readable web page with
- information about the error, used to provide the client developer with additional information about the error.
- Values for the "error\_uri" parameter MUST conform to the URI-reference syntax and thus MUST NOT include characters
- outside the set %x21 / %x23-5B / %x5D-7E.
- state
- REQUIRED if a "state" parameter was present in the client authorization request. The exact value received from the
- client

# /token

- The token request endpoint must support post with Content-Type: application/x-www-form-urlencoded
- the token request endpoint must authenticate the client making the request
- the token request end point must support basic authentication
- the token endpoint must ensure that the authorization code was issued to this client\_ID
- the token endpoint must ensure that the authorization code is valid.
- the token endpoint must ensure that the authorization code is used ONLY once.
- authorization code must expire in a few seconds
- The token endpoint must set Cache-Control: no-store, Pragma: no-cache headers
- The token request endpoint supports the following parameters:
  - grant\_type with value set to authorization\_code, client\_credentials, password, or refresh\_token
  - code with its value set to the authorization code
  - redirect\_URI with its value set to the redirect URI that was provided in the request for the authorization code
  - client\_id; this value is required if the client is not authenticating with the authorization server
- The return payload must include the following:  
access\_token, token\_type, expires\_in refresh\_token, and any other key value pairs.



# Exchange an authorization code for a token

- POST /token HTTP/1.1  
Host: [server.example.com](https://server.example.com)  
Authorization: Basic czZCaGRSa3F0MzpnWDFmQmF0M2JW  
Content-Type: application/x-www-form-urlencoded
- grant\_type=authorization\_code&code=SpIxlOBzQQYbYS6WxSbIA  
&redirect\_uri=https%3A%2F%2Fclient%2Eexample%2Ecom%2Fcb
- HTTP/1.1 200 OK  
Content-Type: application/json; charset=UTF-8  
Cache-Control: no-store  
Pragma: no-cache

```
{  
  "access_token": "2YotnFZFEjr1zCsicMWpAA", "token_type": "Bearer", "expires_in": 3600,  
  "refresh_token": "tGzv3JOkF0XG5Qx2TlKWIA",  
  "example_parameter": "example_value"  
}
```

# Refreshing the access token

- POST /token HTTP/1.1  
Host: [server.example.com](http://server.example.com)  
Authorization: Basic czZCaGRSa3F0MzpnWDFmQmF0M2JW  
Content-Type: application/x-www-form-urlencoded
- grant\_type=refresh\_token&refresh\_token=tGzv3JOkF0XG5Qx2  
TIKWIA

# Methodology for securing rest API

- Client app registers with OAuth/Authorization Server
- Client app request a token
- OAuth provider generates an access token to client APP
- Client app includes access token in every HTTP request using Authorization header
- Client app sets the Authorization header to Bearer {access token}
- The rest API validates the access token
  - What does it need to validate the token?

# Token Validation by Resource Server

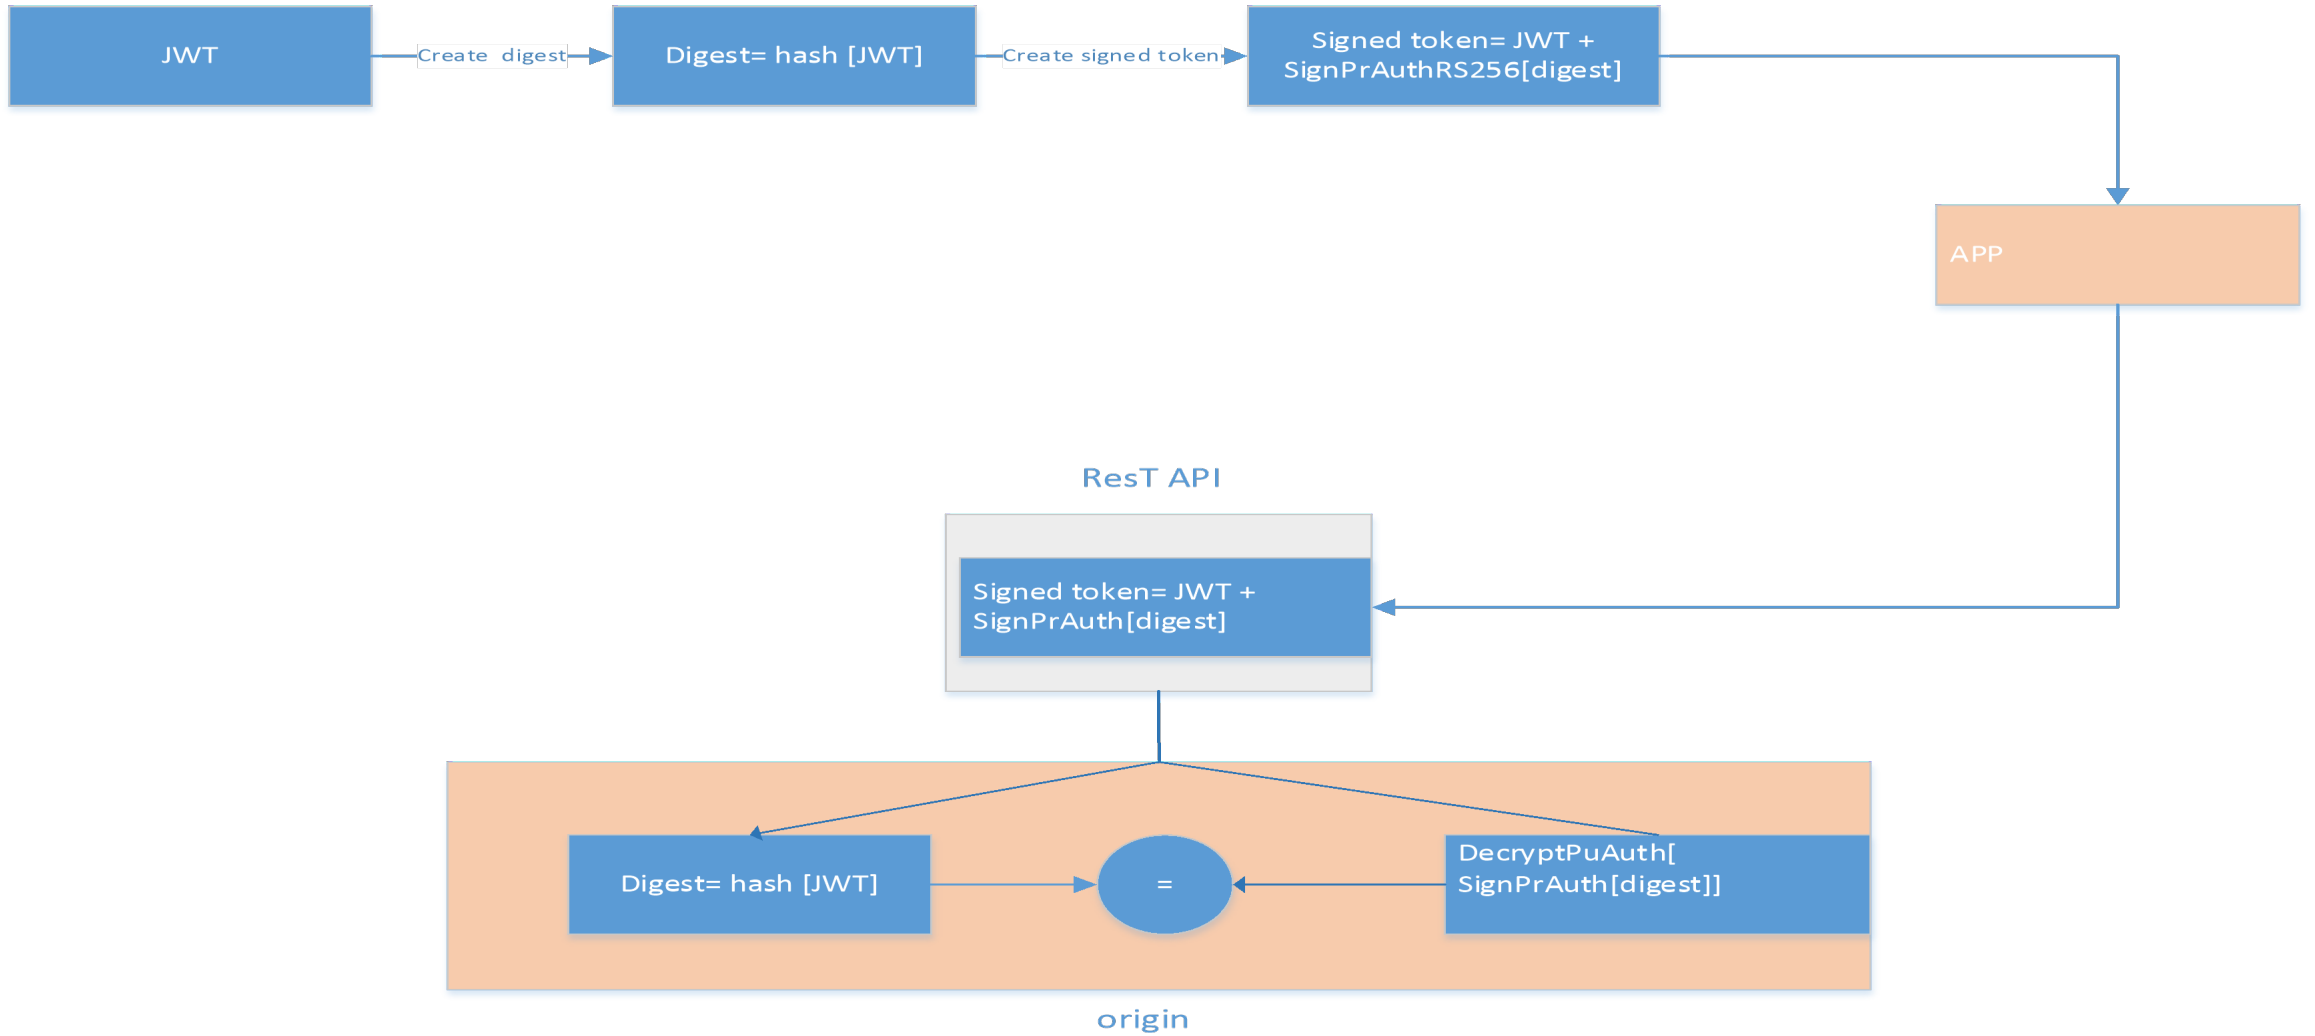
- 1. Validate the structure of a JWT
- 2. Create an “allow list” that contains valid values for iss claims
- 3. Base64decode JWT header, payload
- 4. Retrieve alg and kid from Header
- 5. Retrieve iss from payload
- 6. Compare the value of iss to that stored in the “allowed list”
  - 5. If iss value in allow list, use JWKS to retrieve public key. Otherwise, signature invalid
  - 6. Verify signature
  - 7. Validate any other claims such as scope, aud, exp, etc.

# JWT example

```
{  
  "typ": "JWT"  
}
```

```
{  
  "app": "TEST",  
  "acc": "7888-a9a0-4de2-be72-57775575",  
  "iss": "yyy",  
  "scope": ["read", "write"],  
  "exp": 1561939073,  
  "jti": "jhhhjhg-6cab-lkjjll-8512-kjkkjk",  
  "aud": "/plan/{id}"  
}
```

```
RSASHA256( base64UrlEncode(header) + "." + base64UrlEncode(payload)
```



# Key Distribution

- When using RS 256:
  - Generate a public/private key pair
  - Signer uses the private key to sign the token
  - Rest API uses the public key to verify the signature
  - Rest API must have access to the public key
    - JWK : <https://tools.ietf.org/html/rfc7517>

# Key Rotation (Private) and Distribution (Public)

- Using Kid
- {
- "alg": "RS256",
- "typ": "JWT",
- "kid": "2",
- ~~• "jku": <https://myjwks> ;;; Not recommended to include this~~
- }



# References for token signing

- <https://connect2id.com/products/nimbus-jose-jwt/examples/jwt-with-rsa-signature>
- [https://en.wikipedia.org/wiki/JSON\\_Web\\_Token](https://en.wikipedia.org/wiki/JSON_Web_Token)
- <https://tools.ietf.org/html/rfc7519>
- <https://developers.google.com/oauthplayground/>
- <https://developers.google.com/identity/protocols/oauth2/openid-connect>
- <https://console.developers.google.com/apis/credentials?project=vital-invention-306022>
- <https://accounts.google.com/.well-known/openid-configuration>
- JWT.io
- <https://developers.google.com/identity/protocols/oauth2/openid-connect>

See discussions, stats, and author profiles for this publication at:  
<https://www.researchgate.net/publication/315679274>

Query Service for REST APIs

Article · December 2016

CITATIONS  
0

READS  
3,146

1 author:



[Marwan Sabbouh](#)  
Independent Researcher  
24 PUBLICATIONS 128 CITATIONS  
[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



IT apps as SAAS [View project](#)

# Query Service for REST APIs

Marwan Sabbouh

## 1 Introduction

Big Data and Micro Services architectures make heavy use of ReST APIs that are backed by NoSQL databases and full text search engines. Notably missing from these architectures are relational databases due to the requirements of high throughput, high availability, and low latency on CRUD operations. To provide querying capabilities, full text search engines (e.g. Elastic Search) are starting to provide limited join queries capabilities. Additionally, distributed graph query engines (Titan) and graph query languages (e.g. GraphQL), are gaining in popularity. While these technologies are filling the functional gaps, they are also increasing the complexities of big data architectures. In this paper, I describe an approach for querying data that only leverages full text search and NoSQL databases. The approach builds on the conventions specified by data interchange protocols (e.g. DIP, gData, OData) to define a fast graph based indexing technique and an addressing scheme for all entities consumed by the REST API.

## 2 Background

In this section, we define terminology that helps in explaining the indexing technique. The graph based indexing algorithm runs in real time as the data is ingested through the REST API. A graph is comprised of nodes and relationships. Each graph node is uniquely identified through its type and its unique identifier. A source node is connected to a target node through a relation. The indexing algorithms store the generated indices in the NoSQL database, which are used by the query service for query resolution. In this short paper, I will describe the indexing algorithm and query resolution. Subsequently, we define a relation as the link between a source node and a target node. For addressing, we will assume an addressing scheme like JSON Path. A JSON Object is a collection of property value pairs. There are two types of properties in a JSON Object: simple properties, and object properties. A simple property is a property whose value is not a JSON Object. An object property is a property whose value is a JSON Object. A JSON Object consisting only of simple properties is a node in the graph. An object property is a relation in the graph.

Figure 1 illustrates a JSON document describing comments on a blog entry. The JSON document consists of three JSON Objects. The root object with simple properties `_type`, `_id`, `title` and with the object property `comments`. The value of the `comments` property is another JSON Object with simple properties `_type`, `_id`, `created` and with the object property `author`. The value of the `author` property is another JSON Object with simple properties `_type`, `_id`, and `name`.

```

{
  "_type": "blog",
  "_id": "123456",
  "title": "novel indexing techniques",
  "comments": [ {
    "_type": "comment",
    "_id": "78910",
    "created": "05-12-2017",
    "author": {
      "name": "Michael Smith", // NOTE: Duplicated
      "field": "121314",
      "_type": "person",
      "_inferred": [ "blog-person" ]
    }
  } ]
}

```

Figure 1: Sample JSON Document

Figure 2 shows the graph model of the JSON document.



Figure 2: Graph Representation of Figure 1

### 3 Indexing Algorithm

1. First, the indexing algorithm interprets the nested document shown in figure 1 as a directed graph. This is shown in figure 2.
2. Second, the indexing algorithm makes use of inverseOf inference and transitive inference.  
The resulting graph is shown in figure 3.

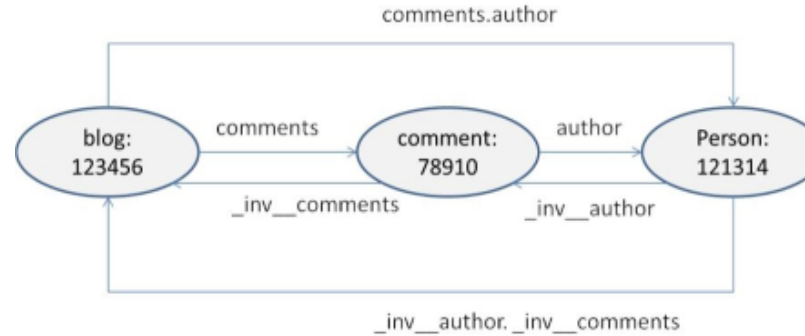


Figure 3: Inferred graph of figure 2

2

Figure 3 shows *\_inv\_comments* as the inverse relationship of *comment*, *\_inv\_author* as

the inverse relationship of *author*, and *\_inv\_author.\_inv\_comments*

as the inverse relationship of *comments.author*. These inverse relationships are generated

automatically by the indexing algorithm. Figure 3 also shows the transitive relationship that

was inferred by the algorithm. In this example, the transitive relationship is shown to be

*comments.author*. Note that, we left the algorithm generate the transitive relationship name and the best it could do is something like "*comments.author*". The algorithm also infers the relationship between nodes *blog:123456* and *person:121314*. This

4. ~~Figure 3 shows~~ the indexing algorithm annotates the target object of all transitive relationships with need of the data modeler defining inverseOf and transitive relationships.

3. The inferred relationship *comments.author* is person\_121314. It contains the property *\_inferred*

the transitive relationship *comments.author* has "blog-person" as its inferred object type. Below in "blog-person".

figure 3.

1. comment\_78910.\_inv\_comments:

[blog\_123456]

2. person\_121314.\_inv\_author: [comment\_78910]

4. person\_121314.\_inv\_author.\_inv\_comments: [blog\_123456]\*

5. blog\_123456\_comments: [comment\_78910]

6. blog\_123456\_comments.author: [person\_121314, blog-person\_121314]\*

7. relation.\_inv\_author.\_inv\_comments\_inverseOf:

[comments.author]\*

```

8. relation__inv__author_inverseOf: [author]*
9. relation_comments.author_inverseOf:
   [_inv__author._inv__comments]*
10. relation_comments.author_inverseOf: inferred blog-person_121314]
*

```

Figure 4: Semantic index

Note that in figure 4 the lines ending with an asterisk indicates inferred knowledge.

#### 4 Querying Algorithm

To illustrate how the indexing algorithm can be used in search, I present the following example.

Suppose

a user would like to find all persons whose first name Michael and who have commented on blogs.

Assuming the availability of an inverted index search engine, e.g. Elastic Search, the query syntax

may look something like this.

To understand the search query, the URI path is of the following form: /search/indexname/{object type}

followed by a typical Lucene query parameters. The fact that the object type is blog, tells the search

engine to return instances of object type blog that matches property name with the value starting with

“Michael”. The JSON path of the property name in the nested document is

“comment.author.name”. In

this case, the JSON document shown in figure 4 will be returned. Please observe that for this query to

remain the same as shown above, /search/indexname/blog?q=comment.author.name: “Michael”.

The query parser does the following.

1. Splits the JSON path “comment.author.name” into relationships comments.author, and person.
2. Find the range of the relation comments.author. This returns type person, and the inferred type

blog-author.

3. Form the query to find all instances of person with the property name starts with Michael: /search/indexname/person?q = name: “Michael\*” & field: “\_id”.

4. Append to the query above the inferred type as a query parameter. The above query becomes:

/search/indexname/person?q = name: “Michael\*” & field: “\_id” & \_inferred: blog-person.

5. Executes the query. It returns the JSON instances of type person with \_inferred property set to

blog-person, and the name property starting with Michael. However, due to the query parameter field: “\_id” the search engine only returns a list of all the JSON instance identifiers

that matches the query. For this simple example, the returned list will contain 121314

7. Find the semantic index the value of the key person\_121314\_

\_inv\_\_author.\_inv\_\_comments “\_id” of the JSON instance of type person.

6. Find the instance of relationship comments.author (found in the first step) by looking up in

the semantic index the value of relation\_comments.author\_inverseOf . This returns relation

\_inv\_\_author.\_inv\_\_comments.

The above algorithm demonstrates how to accomplish join query without data duplication. That is, we could remove the name property from the JSON document of figure 1, and the search query will still return the correct result. This is because the query in step 3 is to search for instances of object type.

## 6 Experimental Results

I implemented the indexing algorithm as part of the ReST API and the query service rest API. I used Redis in conjunction with Elastic Search to implement the system. The ReST API was implemented in Java using Spring Boot. The early results are quite encouraging. We can index a JSON document consisting of 5000 objects in less than 30 ms. A sophisticated join query consisting of three conditions on different nodes in the graph returned in less than 30 ms. Furthermore, the indexing algorithm also proved useful.

## 6 Conclusion

in implementing merge/patch functionality for ReST API, in addition to playing a key role in As we require more functionality from big data technologies, we are faced with increased business rule technological complexities that are preventing many small to medium-size companies from adopting these technologies. For these architectures to become pervasive, techniques that simplify the big data technological stack are critically needed. I believe the approach described here is a step in that direction.



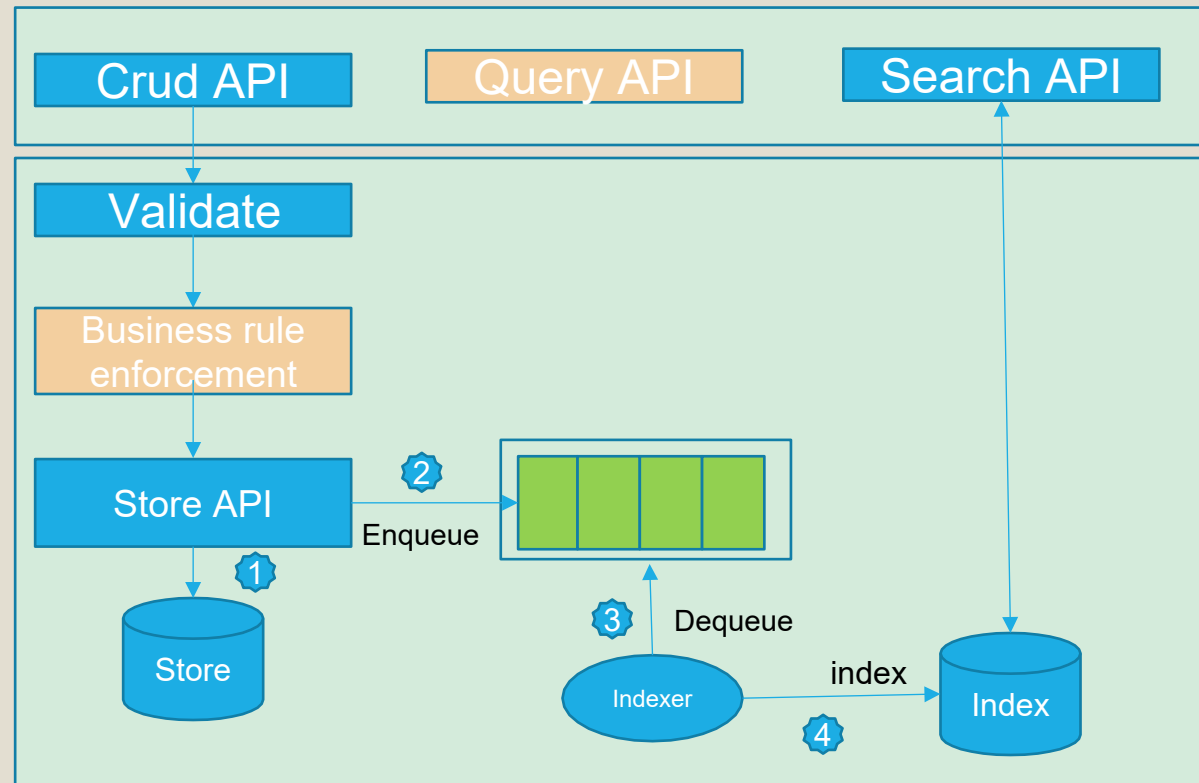
# ADVANCED TOPIC IN BIG DATA



# Quick Review

- By now, you should be familiar with strongly typed data protocols
- You should have reviewed gData, oData, Protocol Buffers
- You should have fair understanding of the overall architecture
- You should have some code working on your laptop

# Architecture



# Prototype Requirements:

Rest API that can handle any structured data in Json

- URIs, status codes, headers, data model, version
- Rest API with support for crd operations
  - Post, Get, Delete
- Rest API with support for validation
  - Json Schema describing the data model for the use case
  - Controller validates incoming payloads against json schema
- The semantics with ReST API operations such as update if not changed/read if changed
  - Update not required
  - Conditional read is required
- Storage of data in key/value store
- Must implement use case provided

# Rest API Specifications



- Data Models
  - Payload structure and serialization
- URI conventions
  - `/ {type} / {id}`
  - `/plan/12xvxc345ssdsds`
- Status Code
  - 200,201
  - 302,304
  - 401, 404, 403, 412, 429
  - 500
- Headers
  - Students should review the HTTP standard headers
  - Various uses of Etag, If-Match, If-None-Match, Authorization in Rest APIs
- Version
  - Accept
  - URL
- Security
- Example: <https://www.hl7.org/fhir/http.html>

# Tooling

- Json simple for Json parsing
- Spring Boot for rest API development
- Elastic Search for search and retrieval capabilities
- Redis for Cache solutions
- Json Schema for schema validation
- Zuul for API Gateway pattern

# But how do I distribute the data?

- single point of failure
- Limited space/storage
- Strongly consistent
- Highly available distributed system
- Seemingly unlimited storage
- What about consistency?



# Key/value stores

- Key readings:

- Dynamo: Amazon's Highly Available Key-value Store :

- <http://www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf>

- Bigtable: A Distributed Storage System for Structured Data:

- <http://static.googleusercontent.com/media/research.google.com/en//archive/bigtable-osdi06.pdf>

-

- CAP Theorem

- Consistency

- Eventual consistency, Read your own write, Strongly consistent

- Availability

- Partition tolerance


- In the presence of network failure, you have to choose between consistency and high-availability

# Problems

- In the presence of many servers, how do I determine the server that stores the object?
  - Consistent hashing to the rescue
- But what if one of the servers fails or the network connection to the server fails?
  - Replication techniques:
    - Primary/backup
    - Active replication
- If I have multiple servers and if an object is stored on more than one server
  - How do I keep the objects consistent?
    - Eventual consistency, strong consistency, weak consistency



# Weak consistency



# Quorum consistency


R=read replica count

W=write replica count

N=replication factor

Q=**QUORUM** ( $Q = N / 2 + 1$ )

- If  $W + R > N$ , you will have consistency
- On read, two of the replica must respond
- On write, two of the replica must make the data durable before acknowledging the right



# Data Modelling

- $K1 \sqsubseteq v1$
- $K2 \sqsubseteq v2$
- $K1?, K2?$

# Consistent hashing

- <http://theory.stanford.edu/~tim/s17/l/l1.pdf>
- How do you map a large number of objects into few servers?
  - $h(x) \bmod n$
- What if the number of servers changes, what would that do to the objects that have already been assigned?
- How do ensure a universal distribution of objects across servers?
- The key idea is:
  - hashing the names of all objects
  - hash the names of all the cache servers  $s$
  - The object and cache names need to be hashed to the same range, such as 32-bit values.

# Key design issues

- Partitioning algorithm
  - Uniform load distribution
- Schema less
- Replication strategy
- Recovering from partial failure
  - Joining a group
    - Load partitioning amongst replicas
- Load rebalancing
- Range query support
- Data versioning
- Support for structured data or simply Blobs
- Marshaling/Unmarshaling
  - How do you store int and floats in redis?

# Mapping of meta-model into key/value store

- JSON payloads can be modeled as a graph.
- [https://www.researchgate.net/publication/315679274\\_Query\\_Service\\_for\\_REST\\_APIs](https://www.researchgate.net/publication/315679274_Query_Service_for_REST_APIs)
- [https://www.researchgate.net/publication/315679444\\_Business\\_Rules\\_for\\_REST\\_APIs](https://www.researchgate.net/publication/315679444_Business_Rules_for_REST_APIs)

How do we map a JSONObject into the key value store?

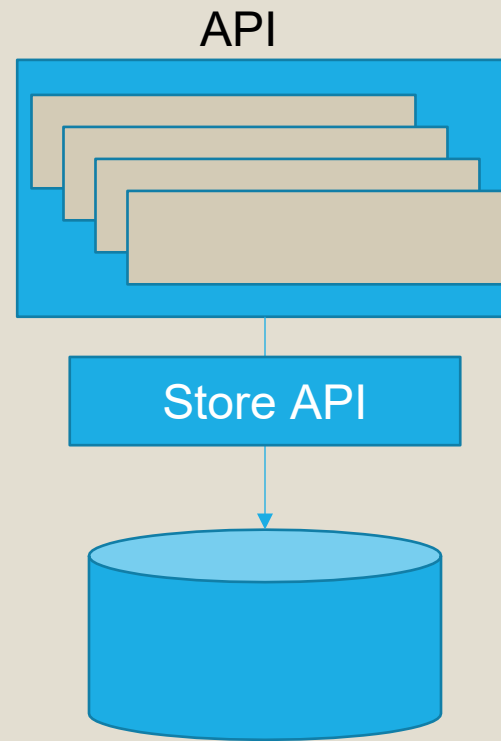
- What is the key signature?
- Do we store the data as a blob?
- Do we store the data as structured?

# Trade-offs between storing the data as a blob versus structured storage

- Storing data as a blob is fast, atomic, reliable
  - But, how do you update the data?
- Storing the data as structured data requires more work on initial creation, but update are much quicker

# A typical design pattern

A compound document  
with nested objects



Should the compound document be decomposed into its constituent objects for storage, and/or indexing, etc...?





# ADVANCED TOPIC IN BIG DATA

# Quick Review

- Syllabus
- JSON
- JSON Schema
- Creating strongly types system with JSON
- The need for validation
- Addressing (briefly)

# Creating strongly typed data with Json

- Every object is an instance of a type
- System exposes aspects, e.g. `_id`, `_type`, etc. that are used in any object.
- Define the type version in the system, and associate with it the property test.
- The property test has datatype array of integers

```
◦ {  
◦     "_name": "version",  
◦     "_type": "Entry",  
◦     "_org": "logoilabs.com",  
◦     "_id": "version",  
  
◦     "properties": [{  
◦         "_name": "test",  
◦         "dataType": ["int"],  
◦         "_type": "Field",  
◦         "_org": "logoilabs.com",  
◦         "_id": "version----test",  
◦         "isFieldOf": "version"}]  
◦ }
```

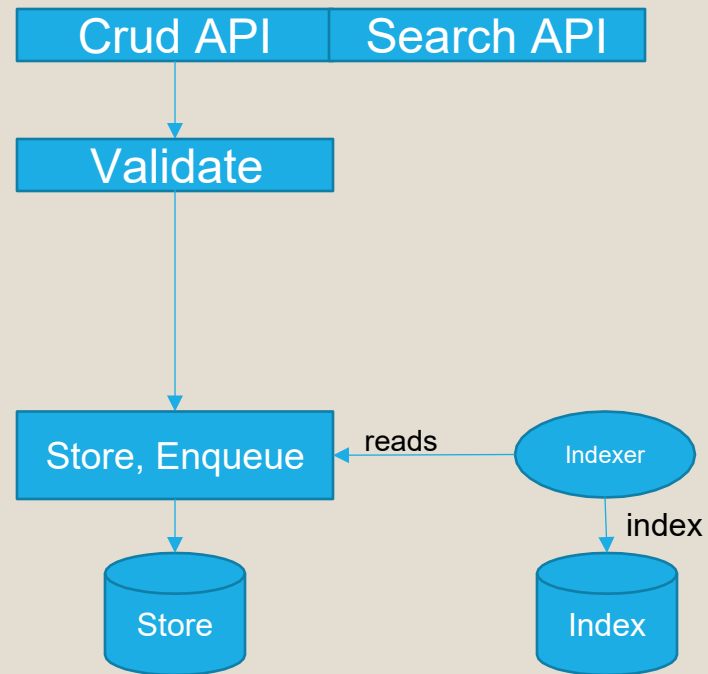
# Variety of strongly typed system

- define types and properties
- define references to objects
- support for inheritance?
- extending the definition of types with additional properties
- system defined types and properties
- aspects
- support for versioning?
- advanced data modeling primitives: intersection, one of, cardinality support, union
- Examples:
  - GDATA, <https://developers.google.com/gdata/>
  - Protobuf, <https://developers.google.com/protocol-buffers/>
  - Microsoft Odata
  - Facebook:GraphQL

# Class exercise

- Come up with a convention to depict a reference to an object in a json payload
- See if you can come up with two different ways of doing it, and then compare the two to choose the better one
- 15 minutes

# Architecture



# Rest API Specifications

- URI conventions
  - /type/id
- Headers
  - Students should review the HTTP standard headers
  - Various uses of Etag, Not Modified Since, Authorization in Rest APIs
- Payload structure and serialization
- Security
- Status Code
- Example: <https://www.hl7.org/fhir/http.html>

- [marwansabbouh@gmail.com](mailto:marwansabbouh@gmail.com)



## Data Sharing for Cloud Computing Platforms

Marwan Sabbouh, Kenneth McCracken, Geoff Cooney  
Social Platform

Here, a Nokia Business  
Cambridge, MA, USA

Marwan.Sabbouh@here.com, Ken.McCracken@here.com,  
Geoff.Cooney@here.com

**Abstract**— Cloud computing platforms consist of a reliable services that are run in the cloud. Typically, consumer applications use software development kits (SDKs) provided by the computing platform services to store, update, and retrieve instances of data in the cloud. Services provided by the cloud computing platform, expose different data modeling paradigms that consumer applications use to interact with the cloud. The service-specific data modeling paradigms and SDKs increase the complexity of data sharing between consumer applications that interact with the different services of the cloud computing platform. To make matters more

composability, flexibility, and scalability. In this paper, we will describe a set of abstractions that can be used to abstract different computing platforms. The database/database [1] [2] [3], in the form of a NoSQL abstractions not only abstract the computing platform, but also NewSQL [4] [5] engine. These distributed enable the data discovery and sharing between applications.

We will further show that these abstractions do not add to the interaction with the cloud system. For example, Amazon Web Services offers DynamoDB [6], a fully managed NoSQL database service that makes it easy to store and retrieve data, for applications requiring varying levels of throughput. DynamoDB data modeling concepts include Tables, Items, and Attributes. Applications wishing to store their data in the cloud, modeling the frequent read and write data attributes as keys, need a system that distributed database contains. However, what if multiple applications need to store and share their data. Certainly, the sharing of the data is enabled through the

type definition, the listing of properties and their constraints, but also contains object-metadata specifying processing directives to the cloud system, property annotations, index definitions (if search queries are desirable for example), and access control policies.

In addition to the above requirements, there is also the requirement that the data model must be processed by the same components of the cloud system that process instance data. Hence, there is a need for the data modeling language to share a common object model with the rest of the instance data.

**XML Schema** defines the requirements are  
the serialization format, which in our case is  
**XML Schema** is a subset of a data interchange protocol providing:  
Object notation (**JSON**) eliminates the potential  
of the Object model stipulates that all data are  
having specific mismatches in the system. **JSON**  
is a type. The object model specifies a unique way  
chosen for many reasons, particularly its simplicity,  
and type exchange. The object model also  
specifies the serialization (9) has been used to  
to extend the remote object. In the context of a  
multiplication (**XML**) a priori the expected of  
the object model may also specify primary keys  
in the modeling language provides the  
structure matches in the system.  
vocabulary for writing data models. Data models can  
serve a  
few purposes in the context of sharing data. First,  
they  
validate the instances. Second, they can be used by  
both  
consumers and providers of data to build a  
common  
understanding for a domain. Third, they contain  
access  
control policies and indexing directives.

5) Universal Resource Identifier (URI) conventions  
URI conventions describe the address of the data in  
the  
cloud system. URI conventions make it possible to  
navigate

the exchange document to extract objects or instances, and a (property, value)-pair.

6) Representational state transfer (REST) APIs for storing and modifying the data

With URI conventions already in place, storing, creating, and retrieving data can easily be accomplished using REST.

II. BACKGROUND  
This paper presents this API and its application to instance data, providing

and data models and advantages and disadvantages. Google's Data

Protocol (gData) [11] is a web protocol for reading, writing, and modifying data on the web. gData supports JSON

serialization. The basic idea is that Google's internal services publish their data using gData, enabling consumers of the

services to consume the data in a uniform way. The gData object model is that of an RSS/Atom feed. gData defines

common definitions of certain objects, called Kinds in gData, but stops short of defining and using a data modeling

language. Therefore, gData is an appropriate paradigm for publishing applications data, but may not be suited as a

protocol for abstracting a NoSQL/NewSQL distributed database.

Open Data Protocol (ODa) [12] is a web protocol that is backed by Microsoft. This protocol is used for creating

and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

As a NoSQL database, HBase has a different paradigm for storing and updating data using web technologies such as Atom feeds and HTTP.

quickly discovered that JSON Schema was not expressive enough to describe the indexing requirements and some of the access control policies that need to be specified.

That is, while we were able to validate instances of data models using JSON Schema, we were unable to validate the data models themselves using JSON Schema, as the data models

contained index definitions for search, and access control policies, whose support required disjoint property/object, among other features missing from JSON Schema.

Since, JSON Schema and our data representation in JSON did not share a common object model; we were unable to process the data models using the same software components we had in

place for the instance data. This meant that in an earlier implementation, the processing of data models was done manually.

Yet, another option that was considered but ultimately rejected was the use of a different modeling language that

Hence, we have arrived at the conclusion that the expressive enough, e.g. Resource Description Framework (RDF) Protocol is needed.

III. PROTOCOL SPECIFICATION  
Data Interchange Protocol (DIP) defines an object model for RDF/OWL. DIP uses special attributes, e.g. for use on the web. It offers a powerful object model, as well as vocabularies for defining types, properties, and classes.

DIP uses a special attribute \_id to specify restrictions, and many other useful data modeling features. DIP defines URI conventions for addressing DIP objects and any statement that is not known to be true is

properties. DIP data modeling language defines a set of properties. These issues present significant challenges to the adoption of RDF/OWL in cloud-based systems.

Additionally, DIP supports data modeling features such as validation of data. Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

Further, that choice would require validation of data. Further, that choice would require validation of data.

access controls settings are applied on the type being defined.

#### A. DIP Object Model

DIP object model treats every object as instance of some type. DIP object model defines a uniform way for 1) creating types that are unique in the cloud system; 2) instantiating instances of those types; 3) identifying those instances in the cloud; 4) referencing remote instances; 4) addressing the instances and their properties. DIP objects are exchanged using DIP documents. A DIP document is comprised of a JSON list with each member of the list is a JSON object.

Fig. 1.

```
{
  "type": "Entry",
  "id": "Field",
  "_name": "Field",
  "_comment": "base class",
  "extends": null,
  "...": {...}
}
```

Fig. 1. DIP document

DIP defines its object model using system properties. DIP object model reserves the use of all property names beginning with `_`. We refer to property names beginning with

`_` as System Properties. System Properties have special meanings, and they can occur in any object. They are integral

to the system and they are not defined in any data model. DIP has the following System Properties that relate to the object model: `_type`, `_name`, `_id`, `_comment`, `_ref`, `_uri`.

`_type` specifies the type of a DIP object, meaning that all

property (name, value)-pairs in the object that are non-System Properties are associated with the type specified as

the value of `_type`. The value of `_type` is a single value string. Fig. 1 shows that the value of `_type` is "Entry". The value of `_type` is interpreted as a relative URI (see `_uri` description). In DIP, Entry is used to define new types. DIP

objects are required to specify `_type`.

`_name` specifies the name of a DIP object. Since `_name` is

Field, Fig. 1 shows that the type being defined is Field.

The `_type`: "Entry", `_id`: "Field", `_name`:

value of `_name` is a single value string, and is unique in the `_comment`: "base class", ...

context of a DIP document. Hence, Fig. 2 shows a DIP document containing a single logical object comprised of two objects with the same name. Objects are required to

specify `_type` to create a single logical object in DIP

In the DIP object model, Fig. 2 and Fig. 1 are semantically equivalent as they define a single logical object

named Field with its "extends" property set to null.

`_id` when combined with `_name` is typically used to specify the key in a NoSQL database that holds the DIP document as its value. The value of `_id` is a single value string, which represents a unique identifier in the cloud system. `_id` must be present in the first object of a DIP document. When `_id` is not present in DIP objects, its value is used to combine `_id`, `_name`, `_type` in a `_uri` is assigned to be that of the first object.

property (name, value)-pair. `_uri` is used as an alternative

notation for `_id`, `_name`, `_type`. Therefore, Fig. 1 could be

stated as shown in Fig. 3.

```
{
  "_uri": "/Field/Field.Entry",
  "_comment": "base class",
  "extends": null,
  ...
},
{...}
}
```

Fig. 3. DIP document using `_uri`

Fig. 3 shows the same logical object as Fig. 1 and Fig. 2.

The value of `_uri` is a single value string and has the form

`/_id/_name/_type`. Hence, Fig. 3 shows that `_id` is

Field, `_name` is Field, and `_type` is Entry. Fig. 3 also shows

the object `_name` provides type extension. Fig. 3 shows

being defined to a mechanism to referencing and create

objects. The value of `_ref` is a list of string.

This string has the same form as the value of `_uri`. However, the

difference between `_ref` and `_uri` is that `_ref` points to an

object that was defined elsewhere, while `_uri` defines

the `_uri`: "/C544C14C-51F0-0001-5FB6-

`2620134344F9F/Nokia:Company")`

```
{
  "_comment": "instance of Person",
  "name": "Marwayne",
  "_ref": { "_ref": "/C544d14d-51F0-0001-5FB6-2620134344F9F/Nokia:Company" }
},
...
}
```

Fig. 4. Typical use of `_ref`

Fig. 4 shows an instance Marwan of type Person. This instance has a property employer. Its value is a reference to an instance Nokia of type Company. It is defined in DIP

DIP Data can be modeled as a property graph [18]. A DIP object is comprised of nodes and edges. A node is characterized by either the value of `_uri`, or the combined value of `_name`, and `_id`. Typically, such node has outgoing edges from it. The properties that are part of the DIP object are either attributes of a node, or outgoing edges from a node. When the property value is another DIP object containing `System Property_ref`, the value of

characterizes a node with incoming edge. That is, when the value of the property is interpreted as a URI, it is a node in the DIP type system supports the basic types graph. Fig. 5, shows the representation of Fig. 4 structures defined by JSON, e.g. array, object, properly graph.

string, true, false, null. That is, a DIP object is a valid JSON object. Entry is the type that all other types use. Entry is used to define Entry types. The type Field is used to define Field types and to associate them with a type. To specify that a type is a primitive type, use the type X in the set the DIP object is a primitive type of type X is sufficient to set the

type property to X in the DIP object, document specifies Entry to be a type. The second object in the DIP document is an instance of type Field. This instance contains the usual System Properties, indicating that the instance name is extends. In addition to the System Properties, it contains other properties: dataType, isFieldOf, and scope. Its autonomy, Fig 6 says that Entry specifies type, it

Instant property called `extends` has a value of `extends` is a type in `extends` the sentence number of `extends` is a field. Furthermore, `extends` is a field of `Entry`. DIP adopts the convention of grouping all properties belonging to type in one DIP document.

```
{
  "_id": "Entry", "_type": "Entry",
  "_name": "Entry",
  "_comment": "base class"
}{
  "_type": "Field", "_name": "extends",
  "dataType": [ "string" ],
  "isFieldOf": "Entry", "sequence": 1
}
```

Fig. 6. Entry definition

2) *Type Field*

Fig. 7 shows excerpts from the definition of Field.

```
{
  "id": "Field", "type": "Entry",
  "name": "Field",
}{
  "type": "Field", "name": "dataType",
  "dataType": "any", "minCardinality": 1,
  "isFieldOf": "Field", "sequence": 1
}{
  "type": "Field", "name": "isFieldOf",
  "dataType": "string",
  "isFieldOf": "Field", "sequence": 2
}{
  "type": "Field", "name": "sequence",
  "dataType": "int",
  "isFieldOf": "Field", "sequence": 13
}{...}]
```

Fig. 7 Field definition

[illegible]

if p has its minCardinality attribute set to a number v, this means every instance of T must contain at least v occurrences of that property. Otherwise, the instance is ~~invalid~~ **invalid** means every instance of T must contain at least v occurrences of that property. Otherwise, the instance is not valid.

value `MultiType` as an instance of `Value` string. However, DIP accomplishes this feat by enclosing two different objects

```
having the same value for name, and different values
for "_id": "C5254000....BD20F34E142C",
"_type": "Person",
"name": "Marwan",
"gender": "male"
},{
  "_type": "Worker",
  "name": "Marwan",
  "employer": "Nokia"
}]
```

Fig. 8 shows an instance named Marwan having two

approaches. Object-oriented multi-typing are almost guaranteed to be more complicated than the standard one properties, intraprocess prefixes (increasingly that the introduction of the properties and behaviors on objects, from the properties belonging to another type.

When a user defines a type T to be an instance of T, ~~Entity~~ **Entity** also be an instance of T2. Otherwise, the ~~Kind of Object;~~ **Kind of Object;** document is not valid.

Figure 1 shows a typical instance of type `Object`. Other figures show the effect of some of the other options available. The `Object` is not valid if it contains an index, i.e. an instance of `SCBEIndex`, the instance of index must also be an instance of `SCBEAccessControl`. That is, if the document defines an index, it must also define access control on the index.

```
{
  "_type": "Object"
  "_name": "SCBEIndex"
  "intersectWith": ["sobeAccessControl"]
}, {
  "_type": "Entry",
  "_name": "SCBEIndex"
}
```

### 5) Extension Mechanism

- Any instance of type A must also be an instance

type B. This is accomplished using the multi-  
typing. Therefore, any document that contains an  
serialization described in section C.3.

the `type` of `obj`. All property and object constraints defined on `type` are applied to `obj` during serialization of that instance of type `B`.

A are applied to the instance of type A.

6. \_\_\_\_\_

ODP serialization offers a straightforward approach:

expressing multiple occurrence of a property. Fig. 10 shows an example.

```

{
  "id": "Person",
  "_type": "Entry",
  "name": "Person"
}, {
  "type": "Field",
  "name": "address",
  "dataType": "string",
  "isFieldOf": "Person",
  "sequence": 1
}, {
  "type": "Field",
  "name": "address ",
  "isFieldOf": "Person",
  "dataType": " ./:Address"
}
}

```

Fig. 10. Definition of property address

Fig. 10 shows that the property dataType occurs twice to state that the property address can have a value either of type *Property Restrictions* string or of type Address.

Earlier in the paper, we have seen examples on how to specify an object property, or property whose value is an instance of a certain type. Sometimes, it is useful to further restrict the value of an object property to not only state that its values are instances of a certain type, but also to specify a restriction on the property belonging to that type. DIP makes that possible, by manipulating the format used for the uri. Recall that the format of uri is of the form "/id/name space". To specify any instance of type Field in the following space, we write the following: " ./:Field;\_op= and; isFieldOf = systemProperties; dataType = string". To specify an instance of type Field by name, we write the following: " ./:name: Field".

To specify any instance of type Field, but with the restriction of having dataType to be string and isFieldOf to be systemProperties, we write the following: " ./:Field; \_op= and; isFieldOf = systemProperties; dataType = string"

To specify any instance of type Field, but with the restriction of having dataType to be string or isFieldOf to be systemProperties, we write the following: " ./:Field; \_op= or; isFieldOf = systemProperties; dataType = string". The above notations for property restrictions have several applications particularly in data modeling languages. For example, DIP data modeling language wishes to provide subclassing functionalities based on restriction of properties. Fig. 11 shows how this can be done using DIP.

Fig.11. Example use of property restriction

Fig. 11 shows the definition of type stringField which is the class of all instances of Field, with dataType = string. Suppose the existence of an instance of Field, but with dataType = string, this would not be a member of the class of all instances of Field. The limitation of cloud systems to offer indexing and search solutions. In this case, applications' data stored in the cloud are indexed in order to provide search functionality. To implement this functionality, cloud systems often allow applications to define their own index definition file. In this approach, each application would have its own index definition. DIP object model makes it possible to index all applications' data using a single index definition. The Table I Index definitions for DIP documents

table	_type	_name	property	value	_ref	simpleType
Header	(string)	(string)	(string)	(string)		

document. Other variation of this header which includes other columns is also possible provided that it maintains the Table II shows the necessary indexed fields that comprise an index definition. The fields tagged as key form a compound key indicating a unique row in the table. The value field is tagged multivalued indicating that field's value is a list. The column labeled simpleType contains the type of the value when it is a simple type, e.g string, int, etc. The column labeled \_ref is an indicator that the value is a pointer

to another object. Table II partially shows the index data for Fig. 10.

Table II Index data of Fig. 10

Person	Entry	address	dataType	["string"]		str
Person	Entry	address	dataType	["Address"]		str

#### D. REST API for Create, Read, Update, and Delete

## Operations

Web data protocols offer their consumers a REST API to create, read, update, and delete (CRUD) their data. DIP offers a fully functional REST API. In this paper, we provide a short summary of these operations in Table III.

From a conceptual perspective, DIP permits clients to create/get/delete a DIP document, and to update/insert/delete a request body for all HTTP operations. For each DIP document. In addition to working at the document and object granularities, DIP permits clients to response body for all HTTP Get, is also a get/update/insert/delete any properties in the DIP document. Therefore, whether retrieving a few additional observations to make on the API.

- Object or the entire document, the structure of also resolve to URI `/_id/_name:T` if it extends `Document`.  
 reply is always the same.  
 so on. In that case, the returned DIP document.
- When receiving a HTTP Post/Put request on `URI`, contains the union of the operations on the two `URI`.  
`URI`.name:T, API enforces that the instance is also `URI`.
- When receiving a HTTP delete request on `URI`, `URI` extends `T`.  
`/_id/_name:T`, API also executes delete operation.

It is an error to create the same DIP document more than once. That is, once the document is created it can only be modified or deleted in subsequent operations.

URL convention	HTTP verb	description
/_id	Post/Delete a Get	Create/Remove/Get DIP document
/_id/_name	Get/Post/ Delete	Retrieve/update/Delete the named instance
/_id/_name:type	Get/Post/ Delete	Retrieve/update/Delete the named/typed instance

#### IV. PERFORMANCE EVALUATION

For DIP to be successful, it must not add significantly to the latency of the NoSQL database. The majority of the latency introduced by DIP comes from the validation of DIP instances. The instance validation is comprised of 1) validate property disjointness ; 2) validate property intersection; 3) check property data types; 4) check for required properties; and 5) check property values and enumeration constraints. It is important to know that this validation time is a function of the number of properties in the object, and not the size of the instance. The tests were run on a personal computer running in Windows 7, with CPU I7-2640M, and 8GB of memory. Most of the overhead imposed by DIP takes place on update operations and not on get. Of that overhead, most of it is spent on validation. We observed that the validation time represents about 6% of our average update operation latency for payloads with less than 200 properties. For large sized payloads e.g. the number of properties approaching 1000, the validation time becomes significant. Therefore, these numbers used in this paper. In one of the project, we used DIP clearly demonstrates that validation can happen synchronously, express constraints related to cloud-based social platform. We found DIP to be very expressive, allowing us to size. However, for larger payload size validation should be asynchronous. After a request has been acknowledged by the server, e.g. by having the more related to the handling of database data structure layer. For distributed databases, DIP provide more scalability with the continuous with the validation. Furthermore, we have available for migrate from low when the data of these numbers can further be improved by utilizing more efficient algorithms for validation. This can be without any impact on the application property, disjoint object of another property and intersection object proved very useful.





# Introduction to big data architecture

# Why big data

- Volume
- Variety
- Velocity
- Extensibility (Schemaless)

# Use case

- A company wishes to provide its employees medical coverage. So, they create medical plans tailored to the employees needs. Each plan consists of large number of covered services, e.g. acupuncture, physical, well-baby visits, emergency room visits, and so on. Additionally, each plan specifies the cost associated with that plan. For example, the co-pay for the various visits, and any deductible that should be met before the patient is reimbursed.
- the company has created a website for its employees where they can view each medical plan and the covered services associated with the plans. Additionally, the website is also used by the plan administrators to create new plans and modify existing plans.
- Is this a use case for big data?

# How can we tell?

- Start by asking a few questions:  
what is the data size of a medical plan?  
What does a medical plan look like? That is,  
how can we model a medical plan?

Other factors?

How many people are viewing the website?  
E.g. throughput rates, latency requirements?

Is there a need to batch import/export plans  
from the system?

# Use case continued

- The analyst responsible for plan creation wants to quickly modify any plans that he created with additional attributes. For example, the analyst may want to remove services and add services to the plan. The analyst may also wish to extend any plan with additional attributes that may not have been foreseen during the design of the system.
- The question is: how can we extend the definition of a plan?

# Technical requirements so far

- Need for data modeling
- Need for CRUD APIs
- Need for batch APIs
- Need for data extensibility
- Need for data validation

# Use case continued

- While an analyst editing a plan, this plan must not be visible to employees.  
Furthermore, other analysts may view, but not edit, this plan
- Hence, the need to secure the system with authentication, and authorization support

# Use case continued

- An employee using the system may find the medical plan that best fit his/her needs by using the search box. A user may search on any attribute
- Technical requirement:
  - need for search



# Assignments

- Json schema: <http://json-schema.org/>
- Json: <HTTP://json.org>
- JSON Parser; JSON Simple
- Jsonpath
- Springboot
- marwansabbouh@gmail.com